

## Module 02 – Data Structure's Effects on Performance


### Foundations Slide Deck

- Searching: most common operation performed by a database system
- Baseline for efficiency is Linear Search
  - ◆ Start at the beginning of a list and proceed element by element until:
    - You find what you're looking for
    - You get to the last element and haven't found it
- Arrays:
  - ◆ Fast for random access BUT slow for random insertions
- Linked list: records are linked through a chain of memory addresses
  - ◆ Slow for random access BUT fast for random insertions
  - ◆ each node stores the data and the address of the next node.
- Binary search

```
def binary_search(arr, target)
    left, right = 0, len(arr) - 1
    while left <= right:
        mid = (left + right) // 2
        if arr[mid] == target:
            return mid
        elif arr[mid] < target:
            left = mid + 1
        else:
            right = mid - 1
    return -1
```

**Binary Search Tree** - a binary tree where every node in the left subtree is less than its parent and every node in the right subtree is greater than its parent.

- Node: entity that contains key or value and pointers to child nodes
- Last nodes of each path are called leaf nodes (they do not contain pointer to child node)
- Edge: link between any two nodes
- Root: topmost node of a tree



- Height: # of edges
- Binary trees can have at most two children

#### Traversal - visit all nodes in a graph

- Inorder: Left → Root → Right    Print after visiting left subtree
- Preorder: Root → Left → Right    Print before visiting subtrees
- Postorder: Left → Right → Root    Print after visiting both subtrees

```
# Binary Tree in Python

class Node:
    def __init__(self, key):
        self.left = None
        self.right = None
        self.val = key

    # Traverse preorder
    def traversePreOrder(self):
        print(self.val, end=' ')
        if self.left:
            self.left.traversePreOrder()
        if self.right:
            self.right.traversePreOrder()

    # Traverse inorder
```

```

def traverseInOrder(self):
    if self.left:
        self.left.traverseInOrder()
    print(self.val, end=' ')
    if self.right:
        self.right.traverseInOrder()

# Traverse postorder
def traversePostOrder(self):
    if self.left:
        self.left.traversePostOrder()
    if self.right:
        self.right.traversePostOrder()
    print(self.val, end=' ')

root = Node(1)

root.left = Node(2)
root.right = Node(3)

root.left.left = Node(4)

print("Pre order Traversal: ", end="")
root.traversePreOrder()
print("\nIn order Traversal: ", end="")
root.traverseInOrder()
print("\nPost order Traversal: ", end="")
root.traversePostOrder()

```

### Searching for a node

- **k** is the key you are searching for.
- **x** is the starting node (root node) from which the search begins.
- **y** is a temporary pointer used to traverse the tree.

### **Explanation of the BST Search Pseudocode:**

The **BST-Search** algorithm searches for a key  $k$  in a binary search tree starting from node  $x$ . Here's how it works, step by step:

BST-Search( $x, k$ )

1:  $y \leftarrow x$

$y \leftarrow x$  means we initialize the pointer  $y$  to point to the root node ( $x$ ). From here, the search will begin.

2: while  $y \neq \text{nil}$  do

The loop *while  $y \neq \text{nil}$  do* ensures that the search continues until the node  $y$  is *nil*, which means either the key has been found or we've reached a leaf node without finding it.

3: if  $\text{key}[y] = k$  then return  $y$

If  $\text{key}[y] = k$ , then the key has been found in the current node  $y$ , and we return this node (return  $y$ ).

4: else if  $\text{key}[y] < k$  then  $y \leftarrow \text{right}[y]$

If  $\text{key}[y] < k$ , the current node's key is smaller than the target key ( $k$ ), so we move to the right child of  $y$ . This follows the property of binary search trees, where the right child contains values greater than the current node.

5: else  $y \leftarrow \text{left}[y]$

If  $\text{key}[y] > k$ , the current node's key is larger than the target key ( $k$ ), so we move to the left child of  $y$ . In a binary search tree, the left child contains values smaller than the current node

6: return ("NOT FOUND")

If  $y$  becomes null, it means we've reached the end of a branch without finding the key  $k$ . In this case, the search concludes by returning "NOT FOUND", indicating that the key is not in the tree.

**Explanation of the BST Insertion Pseudocode:**

BST-Insert(x, z, k)

x is the starting node (usually the root)

z is the new node to insert

k is the key we are inserting into the tree.

1: if x = nil then return "Error"

If x is nil (meaning the tree is empty), return an error message.

2:  $y \leftarrow x$

Initialize y as the current node (start from root, x).

3: while true do{

Start an infinite loop, which will continue until we find the right place for the new node.

4: if  $\text{key}[y] < k$

If the current node's key is smaller than k, move to the right.

5: then  $z \leftarrow \text{left}[y]$

Set z to be the left child of y.

6: else  $z \leftarrow \text{right}[y]$

Otherwise, set z to be the right child of y.

7: if  $z = \text{nil}$  break

If we reach a leaf node (z is nil), break the loop.

8: }

9: if  $\text{key}[y] > k$  then  $\text{left}[y] \leftarrow z$

If y's key is greater than k, insert z as the left child of y.

10: else  $\text{right}[p[y]] \leftarrow z$

Otherwise, insert z as the right child of y.

```
# Insert a node
def insert(node, key):


    # Return a new node if the tree is empty
    if node is None:
        return Node(key)

    # Traverse to the right place and insert the node
    if key < node.key:
        node.left = insert(node.left, key)
    else:
        node.right = insert(node.right, key)

    return node
```

## AVL Tree Rotations Diagram

- Version of Binary search Tree, Self balancing tree
- Goal: Minimize height of the tree and maintain balance factor at each node
  - ◆ Difference must equal -1,0,1 to be balanced



- Balance factor = height |left subtree| - height |right subtree| < 1
  - $X > 1$  means left is greater than right
- Find node of imbalance and rebalance based off of four cases of imbalance

### → Four Cases of Imbalance

#### ◆ Left left insertion

- Node of imbalance z, three nodes leaning left

#### ◆ Left Right insertion

- Node of imbalance z, left then right < (left child node is right heavy)

#### ◆ Right right insertion

- Three nodes leaning right

#### ◆ Right left insertion

- Node of imbalance z, right then left >

### → Rebalancing the cases

- ◆ LL: use a single right rotation

Unbalanced:



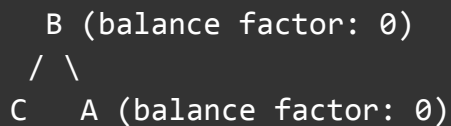
```
def right_rotate(A):
    B = A.left # B is the left child of A
    T2 = B.right # T2 is the right subtree of B

    # Perform rotation
    B.right = A
    A.left = T2

    # Update heights (not shown here)
    update_height(A)
    update_height(B)


    return B # Return the new root of the subtree
```

Balanced:



◆ LR: use a left rotation on left child node, then right rotation on unbalanced node

Left rotation:



```

Right rotation
  A (balance factor: -2)
  /
  C (balance factor: 0)
  /
  B (balance factor: 0)

```

```

    C (balance factor: 0)
   / \
  B  A (balance factor: 0)

```

```

def left_rotate(A):
    B = A.right # B is the right child of A
    T2 = B.left # T2 is the left subtree of B

    # Perform rotation
    B.left = A
    A.right = T2

    # Update heights (not shown here)
    update_height(A)
    update_height(B)

    return B # Return the new root of the subtree

```

◆ RR: use a single left rotation

Unbalanced:

```

  A
  \
   B
   \
    C

```

Balanced:



```

    B (balance factor: 0)
   / \
  A   C (balance factor: 0)

```

- ◆ RL: use right rotation on right child node, then left rotation on unbalanced node

```

Unbalanced
  A (balance factor: +2)
   \
    B (balance factor: -1)
   /
  C (balance factor: 0)
Right Rotation
  A (balance factor: +2)
   \
    C (balance factor: 0)
     \
      B (balance factor: 0)

Left Rotation
  C (balance factor: 0)
 / \
A   B (balance factor: 0)

def right_left_rotate(A):
    # Step 1: Right rotation on the right child of A
    A.right = right_rotate(A.right)

    # Step 2: Left rotation on A
    return left_rotate(A)

```

AVL Node Insertion:

```

class TreeNode:
    def __init__(self, key):
        self.key = key
        self.left = None
        self.right = None
        self.height = 1 # Height of the node (leaf nodes have height
1)

def get_height(node):
    if not node:
        return 0
    return node.height

def get_balance_factor(node):
    if not node:
        return 0
    return get_height(node.left) - get_height(node.right)

def update_height(node):
    if not node:
        return
    node.height = 1 + max(get_height(node.left),
get_height(node.right))

def left_rotate(z):
    y = z.right
    T2 = y.left

    # Perform rotation
    y.left = z
    z.right = T2

    # Update heights
    update_height(z)
    update_height(y)

    return y # New root

def right_rotate(z):

```

```

y = z.left
T3 = y.right

# Perform rotation
y.right = z
z.left = T3

# Update heights
update_height(z)
update_height(y)

return y # New root

def insert(root, key):
    # Step 1: Perform standard BST insertion
    if not root:
        return TreeNode(key)

    if key < root.key:
        root.left = insert(root.left, key)
    else:
        root.right = insert(root.right, key)

    # Step 2: Update the height of the current node
    update_height(root)

    # Step 3: Check the balance factor and rebalance if necessary
    balance = get_balance_factor(root)

    # Left Heavy (balance factor > 1)
    if balance > 1:
        if key < root.left.key: # Left-Left Case
            return right_rotate(root)
        else: # Left-Right Case
            root.left = left_rotate(root.left)
            return right_rotate(root)

    # Right Heavy (balance factor < -1)
    if balance < -1:

```


```

    if key > root.right.key: # Right-Right Case
        return left_rotate(root)
    else: # Right-Left Case
        root.right = right_rotate(root.right)
        return left_rotate(root)

return root # Return the (possibly new) root of the subtree

```

## B Tree



- Root node has two keys: 20,40
  - ◆ Root node has three children - left child contains keys less than 10, middle child contains keys between 20 and 40, right child contains keys greater than 40
- Binary Trees can store many keys in a single node and have multiple child nodes -> decreases height to allow for faster disk access

## B Plus Trees

- Minimize disk access, faster than a disk read,
- $M$  = Maximum # of **keys** in each node (3)
- $M + 1$  = Maximum # of **children** of each node (4)
- Insertion: [04-B+Tree Walkthrough](#)

## Difference between B vs B plus Tree

- B Tree:
  - ◆ Nodes store keys **and** values
- B Plus Tree:

- ◆ Internal nodes only store keys - no values
- ◆ Leaf nodes store keys and values, and they are all link together in a linked list (more effective for range queries)
- ◆ Disk-based storage (like databases)

## Hash Map

- Hash maps store key-value pairs
- Uses **hash function** to map key to an index in an array
  - ◆  $\text{hash}(\text{hello}) = 0 \rightarrow$  put value at index 0
- Collision: when two keys map to same bucket
  - ◆ Resize hashtable to avoid collisions BUT uses a lot more memory
- Collision handling = chaining (storing multiple key-value pairs in the same bucket using a linked list)