

## Module 3 – Moving Beyond the Relational Model

- Benefits of the relational model
  - Relational databases use SQL which is a widely adopted standard
  - ACID compliance ensures transactions are processed reliably
  - Relational models are well-suited for structured data with clear schemas
  - Can handle large data volumes (scalability)
  - There is extensive tooling and expertise available
- Ways a relational database can increase efficiency:
  - Indexing: Speeds up searches by organizing data.
  - Storage Control: Optimizes data layout for efficiency.
  - Column vs. Row Storage: Column-oriented storage benefits analytics; row-oriented storage benefits transactional systems.
  - Query Optimization: Uses execution plans to improve query speed.
  - Caching & Prefetching: Reduces repeated database calls.
  - Materialized Views: Stores query results for fast access.
  - Precompiled Stored Procedures: Speeds up frequently used operations.
  - Data Replication & Partitioning: Improves availability and load balancing.
- Transaction processing – a transaction is a sequence of one or more of the CRUD operations performed as a single, logical unit of work
  - Commit – the entire sequence succeeds (save all changes if successful)
  - Rollback / Abort – the entire sequence fails (revert changes if failure occurs)
  - Transactions help ensure:
    - Data integrity – prevents incomplete updates
    - Error recovery – enables system restoration
    - Concurrency control – manages simultaneous transactions
    - Reliable data storage – ensures durability
    - Simplified error handling – manifest failures systematically
- ACID properties
  - Atomicity – a transaction is all or nothing, it either completes fully or doesn't happen at all
  - Consistency – ensures database transitions from one valid state to another
    - Ex: if a transaction transfers money, it must debit one account and credit another without violating integrity constraints
  - Isolation – transactions execute independently, preventing interference
    - Three common issues arise without proper isolation:
      - Dirty Read – a transaction reads uncommitted changes from another transaction
      - Non-Repeatable Read – a second read during the same transaction returns different data because another transaction modified it

- Phantom Read – a transaction sees new rows added by another transaction before it commits.
  - Durability – once a transaction is committed, it remains so, even after a system failure
- Example of a transaction (SQL Procedure) **check the code in slides 10-11**
  - Implement a money transfer – begin transaction, deduct from sender and credit receiver, check for sufficient funds, rolls back if insufficient funds exist, commits changes if all checks pass
- Challenges of a relational database
  - Schema evolution – sometimes schemas evolve over time and can be difficult to manage
  - Not all apps may need the full strength of ACID compliance
  - Expensive joins – complex queries can be computationally expensive
  - Semi-structured/unstructured data – JSON, XML, etc., don't fit well into rigid schemas
  - Scaling issues – horizontal scaling (across multiple machines) is harder than vertical scaling (upgrading a single machine)
  - Performance needs – some applications require lower latency than traditional relational databases can offer
- Scalability approaches
  - vertical scaling (scaling up) – increase power of a single machine, which is simpler but has physical and financial limits
  - horizontal scaling (scaling out) – uses multiple machines (distributed systems), which is more complex but allows handling larger loads. There are modern systems that make horizontal scaling less problematic
- Distributed systems – a system of independent computers appearing as one computer
  - Key features of a distributed system: computers operate concurrently, computers can fail independently, and there is no shared global clock (?)
  - Distributed databases
    - Data is stored on multiple nodes, often replicated.
    - MySQL, PostgreSQL support replication/sharding
    - NoSQL databases are often distributed, supporting one or both models
    - Key challenge – network failures are inevitable (network partitioning)
    - Distributed databases can be relational or non-relational
- The CAP theorem — a distributed database cannot guarantee all three simultaneously
  - Consistency – every read gets the most recent write or an error
  - Availability – every request gets a response (even if not the latest)

- Partition tolerance – the system continues functioning despite network failures
- CAP theorem - DB view
  - Consistency\*: Every user of the DB has an identical view of the data at any given instant (\*note that the definition of consistency in CAP is different from that of ACID)
  - Availability: In the event of a failure, the database remains operational
  - Partition Tolerance: The database can maintain operations in the event of the network's failing between two segments of the distributed system
  - Real world database trade-offs
    - CP (Consistency + Partition Tolerance): Prioritizes data correctness but may sacrifice availability during failures, includes MongoDB, HBase, Redis
    - AP (Availability + Partition Tolerance): Always responds but may return outdated data, includes CouchDB, Cassandra, DynamoDB
    - CA (Consistency + Availability): Works only if network partitions never happen (which is unrealistic), includes RDBMS, PostgreSQL, MySQL
  - If a system must always respond, it cannot guarantee consistency if partitions exist, a trade-off is always necessary
- Relational databases offer ACID properties, structured data handling, and strong performance tuning
- Distributed databases help with scalability but face CAP theorem limitations
- Choosing the right database depends on application needs (performance, consistency, scalability)

#### Module 4 – Replicating Data

- Benefits of replicating (distributing) data
  - Scalability & High Throughput – handles growing data volumes or increasing read/write loads beyond a single machine's capacity
  - Fault Tolerance & High Availability – ensures application uptime even if some machines fail
  - Latency Reduction – improves response times by placing data closer to users worldwide
- Challenges of distributed data
  - Consistency issues – updates must propagate across multiple nodes
  - Increased application complexity – applications often need to handle distributed read/ writes
- Scaling strategies
  - Vertical Scaling (Scaling Up)

- Shared memory architecture – a single, centralized server with fault tolerance via hot-swappable components
    - Shared disk architecture – multiple machines share a fast storage system, but write contention limits scalability
  - Horizontal Scaling (Scaling Out)
    - Shared nothing architecture – each node has its own CPU, memory, and disk
    - Nodes communicate via a network and are often geographically distributed
    - Uses commodity hardware (cheaper than scaling up)
- Replication vs. Partitioning
  - Replication duplicates data across multiple nodes
  - Partitioning (sharding): divides data into subsets across different nodes
- Common replication strategies
  - Distributed databases typically use one of these replication models:
    - Single Leader Model (most common)
      - All writes go to a single leader.
      - Leader sends replication info to followers.
      - Clients can read from leader or followers.
    - Multiple Leader Model
      - Multiple nodes can accept writes.
      - Each leader synchronizes changes across the system.
      - More complex, but better for distributed workloads.
    - Leaderless Model
      - No designated leader
      - All nodes can accept reads and writes.
      - Data consistency is maintained via quorum-based reads and writes.
- Leader based replication is most common and widely used in both relational and NoSQL databases
  - Relational: MySQL, PostgreSQL, Oracle, SQL Server.
  - NoSQL: MongoDB, RethinkDB, LinkedIn's Espresso.
  - Messaging Systems: Kafka, RabbitMQ.
- How is replication data transmitted to followers? Methods below
  - Statement-Based Replication
    - Sends SQL statements (INSERT, UPDATE, DELETE) to replicas.
    - Issues: error prone due to non-deterministic functions (NOW()), triggers, and concurrency problems.
  - Write-Ahead Log (WAL) Replication
    - Logs every byte-level change in the database.

- Issue: Requires all nodes to use the same storage engine.
  - Logical (Row-Based) Log Replication
    - Stores row changes (before/after state) in a transaction log.
    - Benefit: Decoupled from storage engine, easier to interpret
  - Trigger-Based Replication
    - Logs changes into a separate table when triggers fire.
    - Benefit: Customizable for application needs.
    - Issue: Can be error-prone
- Synchronous vs. asynchronous replication
  - Synchronous Replication
    - Leader waits for followers to confirm before committing changes.
    - Guarantees strong consistency but increases latency.
  - Asynchronous Replication
    - Leader does not wait for confirmation from followers.
    - Improves availability and performance but can lead to data inconsistencies
- What happens when the leader fails? Challenges below
  - Electing a new leader
    - Use a consensus strategy (e.g., node with the most recent updates).
    - Use a dedicated controller node to assign leadership.
  - Client redirection
    - Update clients to write to the new leader.
  - Handling data loss (if using asynchronous replication):
    - New leader may not have all writes.
    - Options:
      - Recover missing writes.
      - Accept some data loss.
  - Avoiding "Split Brain" scenarios
    - If the old leader recovers, there must be a way to prevent multiple leaders from conflicting.
  - Failure detection
    - Determining the optimal timeout before electing a new leader
- Replication lag – the time it takes for updates on the leader to be reflected on all the followers
  - Synchronous replication: Replication lag causes writes to be slower and the system to be more brittle as num followers increases.
  - Asynchronous replication: We maintain availability but at the cost of delayed or eventual consistency. This delay is called the inconsistency window.
- Read-After-Write Consistency
  - Scenario: A user adds a comment on Reddit but does not see it immediately.

- Solution 1: Always read the most recent data from the leader.
- Solution 2: Temporarily switch to reading from the leader for "recently updated" data.
- Challenge:
  - Followers are geographically close to users, but now we must route reads to a distant leader for consistency.
- Implementation (as in scenario):
  - Method 1: Modifiable data (from the client's perspective) is always read from the leader.
  - Method 2: Dynamically switch to reading from leader for "recently updated" data.
    - For example, have a policy that all requests within one minute of last update come from leader.
- Monotonic read consistency
  - Issue – if a user reads data from different followers, they might see older data after newer data. Anomalies occur when a user reads values out of order from multiple followers.
  - Solution – ensure that once a user reads newer data, they will never read older data. Monotonic read consistency ensures that when a user makes multiple reads, they will not read older data after previously reading newer data.
- Consistent Prefix Reads
  - Issue – when different partitions replicate at different rates, data may be read out of order. There is no global write consistency
  - Solution – Consistent Prefix Reads Guarantee ensures that data is always read in the order it was written.
- Summary
  - Data replication improves scalability, fault tolerance, and latency but introduces consistency challenges.
  - Common replication models:
    - Single Leader (Master-Slave)
    - Multiple Leader
    - Leaderless
  - Replication methods vary in complexity and performance trade-offs.
  - Synchronous vs. Asynchronous Replication:
    - Synchronous ensures consistency but increases latency.
    - Asynchronous is faster but may lead to inconsistencies.
  - Leader failure introduces major challenges, requiring leader election, data reconciliation, and client redirection.

- Replication lag affects read consistency, and various strategies like Read-After-Write and Monotonic Reads help mitigate issues
- Consistent Prefix Reads ensure ordered data retrieval across partitions.

## Module 5 – NoSQL & KV DBs

- Database Concurrency: Handling Multiple Transactions
  - ACID Transactions & Pessimistic Concurrency
    - ACID (Atomicity, Consistency, Isolation, Durability): Ensures transactions are processed reliably.
    - Uses locks to prevent conflicts.
      - Example: Borrowing a library book. If you have it, no one else can until you return it.
  - Optimistic Concurrency
    - Assumes conflicts are rare, so transactions don't lock data.
    - Instead, they check at the end of a transaction whether another process has modified the data.
    - Works well in systems with many reads but few writes, like backups or analytical databases
- Origin of noSQL databases was meant for non-SQL databases (in 1998) but now it means Not Only SQL, which are flexible, scalable databases that don't necessarily use SQL. they are designed to handle unstructured, web-scale data efficiently
- CAP Theorem Review (Trade-offs in Distributed Databases)
  - You can only have two out of three:
    - Consistency: Every user sees the same, latest data.
    - Availability: The system is always up and running.
    - Partition Tolerance: The system keeps working even if some network connections fail.
  - Real-World CAP Trade-offs
    - Consistency + Availability: Always provides fresh data but struggles during network failures.
    - Consistency + Partition Tolerance: Ensures accurate data but may reject requests when failures occur.
    - Availability + Partition Tolerance: Keeps running even during failures but might return slightly outdated data.
- Alternative to ACID: BASE Model (Used in NoSQL) – instead of strict consistency, NoSQL databases use BASE:
  - Basically Available: System remains operational, but some data may be temporarily missing.
  - Soft State: Data updates over time, even without new inputs.

- Eventual Consistency: Data will become consistent eventually, but not always immediately.
- Types of NoSQL Databases
  - Key-Value Stores (Simplest, fastest)
    - Stores data as key = value pairs.
    - Examples: Redis, DynamoDB
  - Document Stores
    - Stores JSON-like documents.
    - Examples: MongoDB, CouchDB
  - Column-Family Stores
    - Optimized for large-scale data storage.
    - Examples: Cassandra, HBase
  - Graph Databases
    - Used for relationships between data points.
    - Examples: Neo4j, ArangoDB
- Key-Value (KV) Databases Explained
  - Simpler than relational databases (No complex tables).
  - Super fast because they use in-memory storage and hash tables.
  - No complex queries—just quick data retrieval by key.
  - Key-Value Store Benefits
    - Simplicity: Easy to use.
    - Speed: Fetching a value by key is super fast ( $O(1)$  time complexity).
    - Scalability: Horizontally scalable—just add more servers
- Common use cases for Key-Value Databases
  - Experiment Data Storage (store test results, machine learning data, store intermediate results from data preprocessing and EDA store experiment or testing (A/B) results w/o prod db).
  - Feature Store (store precomputed ML model features, store frequently accessed feature, low-latency retrieval for model training and prediction)
  - Model monitoring (store key metrics about the performance of the model, for example in real-time inferencing)
  - Session Management (store user sessions for quick access, everything about the current session can be stored via a single PUT or POST and retrieved with a single GET and very fast).
  - User Profiles & Preferences (quick lookup of user settings, User info could be obtained with a single GET operation... language, TZ, product or UI preferences).
  - Shopping Carts (fast storage of items in a cart, Cart data is tied to the user needs to be available across browsers, machines, sessions).



- Caching (speed up web requests by temporarily storing frequently accessed data)
- Introduction to Redis (Popular Key-Value Database) and sometimes called a data structure store
  - Open-source, in-memory database (fast and lightweight).
  - Can be used as:
    - Key-Value store
    - Graph Database
    - Full-text search
    - Time-series database
  - Fast operations: Handles 100,000+ writes per second!
  - Persistence options:
    - Takes snapshots at intervals.
    - Maintains an append-only log for crash recovery
  - Does not handle complex data and no secondary indexes, only supports lookup by key
- Redis Data Types
  - Keys: Usually strings, but can be any binary sequence.
  - Values:
    - Strings (Basic text or binary values).
    - Lists (Ordered collections, useful for queues, chat messages).
    - Sets (Unique, unordered elements).
    - Sorted Sets (Like sets but with a ranking system).
    - Hashes (Field-value pairs, similar to objects).
    - Geospatial Data (Store and query location-based data)
- Setting Up Redis
  - Using Docker
    - Download and run the Redis image.
    - Expose port 6379 (default Redis port).
    - Avoid exposing it publicly for security reasons.
  - Using DataGrip
    - File > New > Data Source > Redis
    - Connect using port 6379.
    - Test the connection.-
- Basic Redis Commands - command and description:
  - SET key value — Store a key-value pair
  - GET key — Retrieve the value for a key
  - DEL key — Delete a key
  - EXISTS key — Check if a key exists
  - EXISTS key — Check if a key exists (returns 1 if exists, 0 otherwise)

- EXPIRE key seconds — Set a key to expire after a set time
- TTL key — Check how much time is left before a key expires
- FLUSHDB — Deletes all keys in the current Redis database
- FLUSHALL — Deletes all keys in all Redis databases
- Example:
  - SET username "JohnDoe"
  - GET username # Output: "JohnDoe"
  - EXPIRE username 60 # Key expires in 60 seconds
  - TTL username # Check remaining time before expiration
- Working with numbers in redis
  - INCR key — Increase an integer value by 1
  - DECR key — Decrease an integer value by 1
  - INCRBY key amount — Increase by a specific amount
  - DECRBY key amount — Decrease by a specific amount
  - Example:
    - SET counter 10
    - INCR counter # counter = 11
    - INCRBY counter 5 # counter = 16
    - DECR counter # counter = 15
    - DECRBY counter 3 # counter = 12
- Working with Hashes (Key-Value Store for Objects) - a value of KV entry is a collection of field-value pairs
  - Use Cases:
    - Can be used to represent basic objects/structures
    - number of field/value pairs per hash is  $2^{32}-1$
    - practical limit: available system resources (e.g. memory)
    - Session information management
    - User/Event tracking (could include TTL)
    - Active Session Tracking (all sessions under one hash key)
  - Basic hash commands:
    - HSET key field value — Set a field inside a hash
    - HGET key field — Retrieve a specific field's value
    - HGETALL key — Get all fields and values in a hash
    - HMGET key field1 field2 — Get multiple field values
    - HINCRBY key field amount — increase a numerical field by a certain amount
    - Example:
      - HSET user:1 name "Alice" age 30 city "Boston"
      - HGET user:1 name # Output: "Alice"

- HGETALL user:1 # Output: name: "Alice", age: 30, city: "Boston"
  - HINCRBY user:1 age 2 # Age is now 32
- Linked lists crash course
  - Sequential data structure of linked nodes (instead of contiguously allocated memory)
  - Each node points to the next element of the list (except the last one - points to nil/null)
  - $O(1)$  to insert new value at front or insert new value at end
- List Data Type (Queues & Stacks) - value of KV Pair is linked lists of string values
  - A List in Redis is like an array or queue. It stores ordered elements
  - Use Cases
    - Implementing message queues (e.g., order processing)
    - queue management & message passing queues (producer/consumer model)
    - logging systems (easy to keep in chronological order)
    - build social media streams/feeds
    - batch processing by queueing up a set of tasks to be executed sequentially at a later time
    - Managing task lists (e.g., batch jobs).
    - Storing chat messages (e.g., live messaging apps).
  - Queue (FIFO) Commands - command and description:
    - LPUSH queue value — Add to the start of the queue
    - RPOP queue — Remove from the end of the queue
  - Stack (LIFO) Commands - command and description
    - LPUSH stack value — Add to the top of the stack
    - LPOP stack — Remove from the top of the stack
  - A List in Redis is like an array or queue. It stores ordered elements.
  - Commands for lists
    - LPUSH key value — Add to the front of the list (left push)
    - RPUSH key value — Add to the end of the list (right push)
    - LPOP key — Remove and return the first element (FIFO)
    - RPOP key — Remove and return the last element (LIFO)
    - LLEN key — Get the length of the list
    - LRANGE key start stop — Get a subset of the list
  - Example:
    - LPUSH todo "Task 1"
    - LPUSH todo "Task 2"
    - RPUSH todo "Task 3"
    - LRANGE todo 0 -1 # Output: ["Task 2", "Task 1", "Task 3"]

- LPOP todo # Output: "Task 2" (removed)
- Sets (Unique Collections)
  - Use Cases
    - Tracking unique website visitors.
    - Managing user groups (e.g., students enrolled in a course).
    - Access control lists (who has permission to see something).
  - Set Commands - command and description
    - SADD key value — Add an item to the set
    - SREM key value — Remove an item from the set
    - SISMEMBER key value — Check if an item exists in the set
    - SCARD key — Get the number of items in the set
    - SMEMBERS key — Get all items in the set
    - Example:
      - SADD students "Alice"
      - SADD students "Bob"
      - SADD students "Charlie"
      - SADD students "Alice" # Will not be added again (duplicate)
      - SMEMBERS students # Output: ["Alice", "Bob", "Charlie"]
      - SISMEMBER students "Bob" # Output: 1 (true)
      - SCARD students # Output: 3 (size of set)
- Sorted sets (ranking system): A Sorted Set is like a regular Set, but each item has a score. This is useful for leaderboards, rankings, and priority queues.
  - Commands for sorted sets
    - ZADD key score value — Add a value with a ranking (score)
    - ZRANGE key start stop WITHSCORES — Get values in order (lowest to highest)
    - ZREVRANGE key start stop WITHSCORES — Get values in reverse order (highest to lowest)
    - ZSCORE key value — Get the score of a specific value
  - Example:
    - ZADD leaderboard 100 "Alice"
    - ZADD leaderboard 200 "Bob"
    - ZADD leaderboard 150 "Charlie"
    - ZRANGE leaderboard 0 -1 WITHSCORES # Output: ["Alice", 100, "Charlie", 150, "Bob", 200]
    - ZREVRANGE leaderboard 0 -1 WITHSCORES # Output: ["Bob", 200, "Charlie", 150, "Alice", 100]
- JSON data in redis allows you to stored structured data
  - Commands for json in redis
    - JSON.SET key . value — Store a JSON object

- JSON.GET key — Retrieve a JSON object
  - JSON.GET key path — Retrieve a specific field in JSON
- Example:
  - JSON.SET user:1 . '{"name": "Alice", "age": 30, "city": "Boston"}'
  - JSON.GET user:1 # Output: {"name": "Alice", "age": 30, "city": "Boston"}
  - JSON.GET user:1 .name # Output: "Alice"
- Summary:
  - Hashes store structured objects like user profiles.
  - Lists work as queues and stacks.
  - Sets store unique items like friend lists.
  - Sorted Sets keep ranked items like leaderboards.
  - JSON support allows for structured storage.

## Module 6 – Redis and Python

- After installation of redis, connect to a redis server with the following code
  - import redis
  - redis\_client = redis.Redis(
    - host='localhost', # Or 127.0.0.1 if using Docker
    - port=6379, # Default Redis port
    - db=2, # Redis has 16 databases (0-15)
    - decode\_responses=True # Converts bytes to strings automatically
  - )
- Basic string operations in redis with python
  - # Set a key-value pair
  - redis\_client.set('clickCount:/abc', 0)
  - 
  - # Get value of a key
  - val = redis\_client.get('clickCount:/abc')
  - 
  - # Increment the value
  - redis\_client.incr('clickCount:/abc')
  - 
  - # Get updated value
  - ret\_val = redis\_client.get('clickCount:/abc')
  - print(f'Click count = {ret\_val}')
- Multiple values at once
  - # Set multiple key-value pairs at once
  - redis\_client.mset({
    - 'key1': 'val1',

- 'key2': 'val2',
- 'key3': 'val3'
- })
- 
- # Get multiple values at once
- print(redis\_client.mget('key1', 'key2', 'key3'))
- # Output: ['val1', 'val2', 'val3']
- List commands in redis with python
  - # Create a list and add elements
  - redis\_client.rpush('names', 'mark', 'sam', 'nick')
  - 
  - # Get all values from the list
  - print(redis\_client.lrange('names', 0, -1))
  - # Output: ['mark', 'sam', 'nick']
- More list commands
  - rpush(key, value1, value2, ...) — Add values to the end of the list
  - lpush(key, value1, value2, ...) — Add values to the start of the list
  - lpop(key) — Remove and return the first element
  - rpop(key) — Remove and return the last element
  - lrange(key, start, stop) — Get values from a list (similar to slicing in Python)
  - llen(key) — Get the number of elements in the list
- Hash commands in redis with python
  - # Store a user profile in a Redis hash
  - redis\_client.hset('user-session:123', mapping={
  - 'first': 'Sam',
  - 'last': 'Uelle',
  - 'company': 'Redis',
  - 'age': 30
  - })
  - 
  - # Retrieve all fields
  - print(redis\_client.hgetall('user-session:123'))
  - # Output: {'first': 'Sam', 'last': 'Uelle', 'company': 'Redis', 'age': '30'}
- More hash commands
  - hset(key, field, value) — Set a field in a hash
  - hget(key, field) — Get a field's value
  - hgetall(key) — Get all fields and values
  - hdel(key, field) — Delete a field
  - hlen(key) — Get the number of fields

- Using redis pipelines for efficiency; normally, sending multiple commands to Redis one at a time can be slow. Pipelines allow batch processing, reducing network overhead and are ideal for bulk inserts or batch operations
  - `r = redis.Redis(decode_responses=True)`
  - `pipe = r.pipeline()`
  - 
  - `# Batch insert multiple keys at once`
  - `for i in range(5):`
  - `pipe.set(f"seat:{i}", f"#{i}")`
  - 
  - `# Execute all at once`
  - `set_5_result = pipe.execute()`
  - `print(set_5_result) # Output: [True, True, True, True, True]`
  - 
  - `# Retrieve multiple keys in a batch`
  - `pipe = r.pipeline()`
  - `get_3_result = pipe.get("seat:0").get("seat:3").get("seat:4").execute()`
  - `print(get_3_result) # Output: ['#0', '#3', '#4']`
- Redis as a feature store for machine learning
  - Example: storing user engagement data for machine learning
    - `redis_client.hset('user:1001', mapping={`
    - `'click_rate': 0.75,`
    - `'session_time': 180, # seconds`
    - `'purchase_likelihood': 0.92`
    - `})`
    - 
    - `# Fast retrieval during ML inference`
    - `user_data = redis_client.hgetall('user:1001')`
    - `print(user_data)`
    - `# Output: {'click_rate': '0.75', 'session_time': '180',`
    - `'purchase_likelihood': '0.92'}`
  - Why use redis for machine learning?
    - Fast Reads/Writes: In-memory storage ensures low latency.
    - Scalability: Redis can store millions of ML feature sets.
    - Simple API: Works well with real-time AI models